

CXL□□□□□□

**CXL**□□□□□□

□□**CXL**□□□□□□

□□**CXL**□□□□

□□**CXL**□□□□□

□□**CXL**□□□

□□**CXL**□□□

□□**CXL**□□□□



CXL

**Compute Express Link (CXL)**   
   
 CXL   
 PCIe   
 PCIe   
 PCIe

CXL 2019 Inter 2019 3 Meta Microsoft HPE CXL

CXL□□□□□□

1. □□□□□□□□□□□□□□

2. □□□□□□□


3. □□□□□□□□□□□


4. □□□□□□□□□□□□□□□□□□





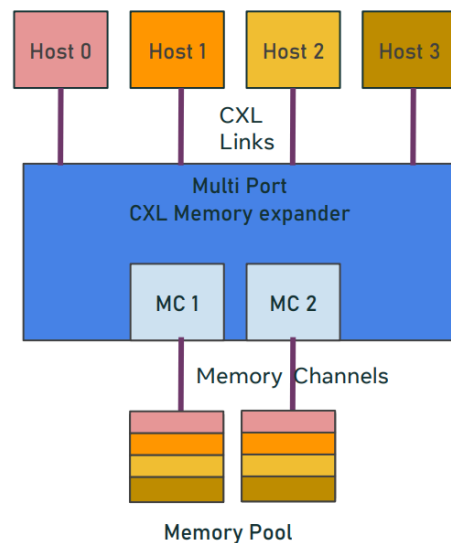
[illegible][illegible]

[illegible]




## Direct attached

Add Capacity  
Add Bandwidth  
Slower-cheaper tier



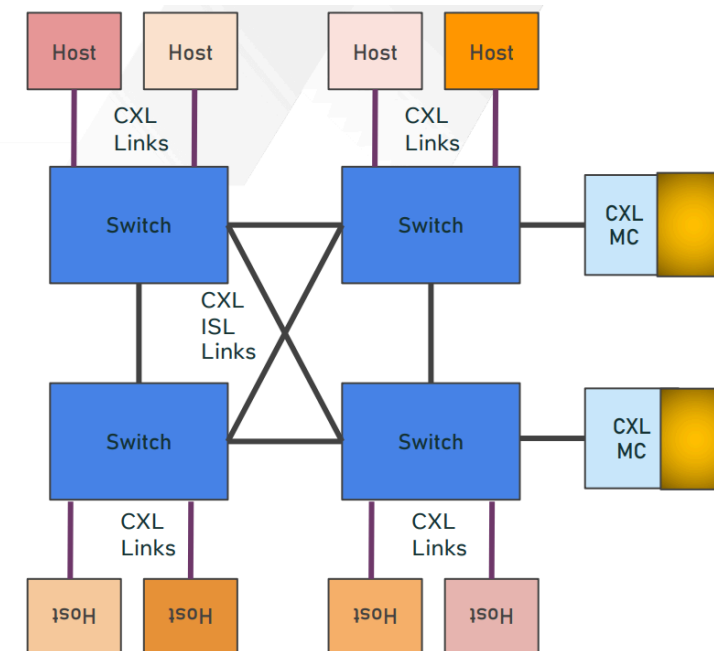
## Pooled Memory

Amortize CXL infra cost  
Flexible allocation



## Shared Memory

Deduplication  
Host2host communication  
large datasets



## Fabric Memory

Scaling to huge datasets

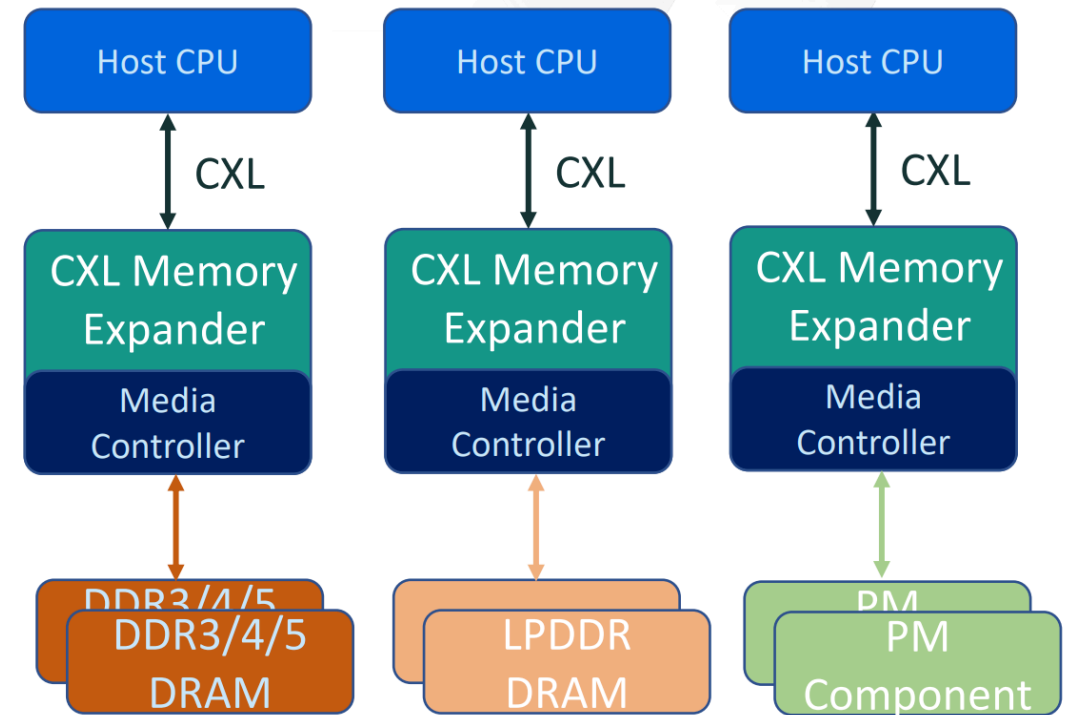


# CXL□□□□□

Features	CXL 1.0 / 1.1	CXL 2.0	CXL 3.0	CXL 3.1
Release date	2019	2020	August 2022	November 2023
Max link rate	32GTs	32GTs	64GTs	64GTs
Flit 68 byte (up to 32 GTs)	✓	✓	✓	✓
Flit 256 byte (up to 64 GTs)			✓	✓
Type 1, Type 2 and Type 3 Devices	✓	✓	✓	✓
Memory Pooling w/ MLDs		✓	✓	✓
Global Persistent Flush		✓	✓	✓
CXL IDE		✓	✓	✓
Switching (Single-level)		✓	✓	✓
Switching (Multi-level)			✓	✓
Direct memory access for peer-to-peer			✓	✓
Enhanced coherency (256 byte flit)			✓	✓
Memory sharing (256 byte flit)			✓	✓
Multiple Type 1/Type 2 devices per root port			✓	✓
Fabric capabilities (256 byte flit)			✓	✓
Fabric Manager API definition for PBR Switch				✓
Host-to-Host communication with Global Integrated Memory (GIM) concept				✓
Trusted-Execution-Environment (TEE) Security Protocol				✓
Memory expander enhancements (up to 32-bit of meta data, RAS capability enhancements)				✓

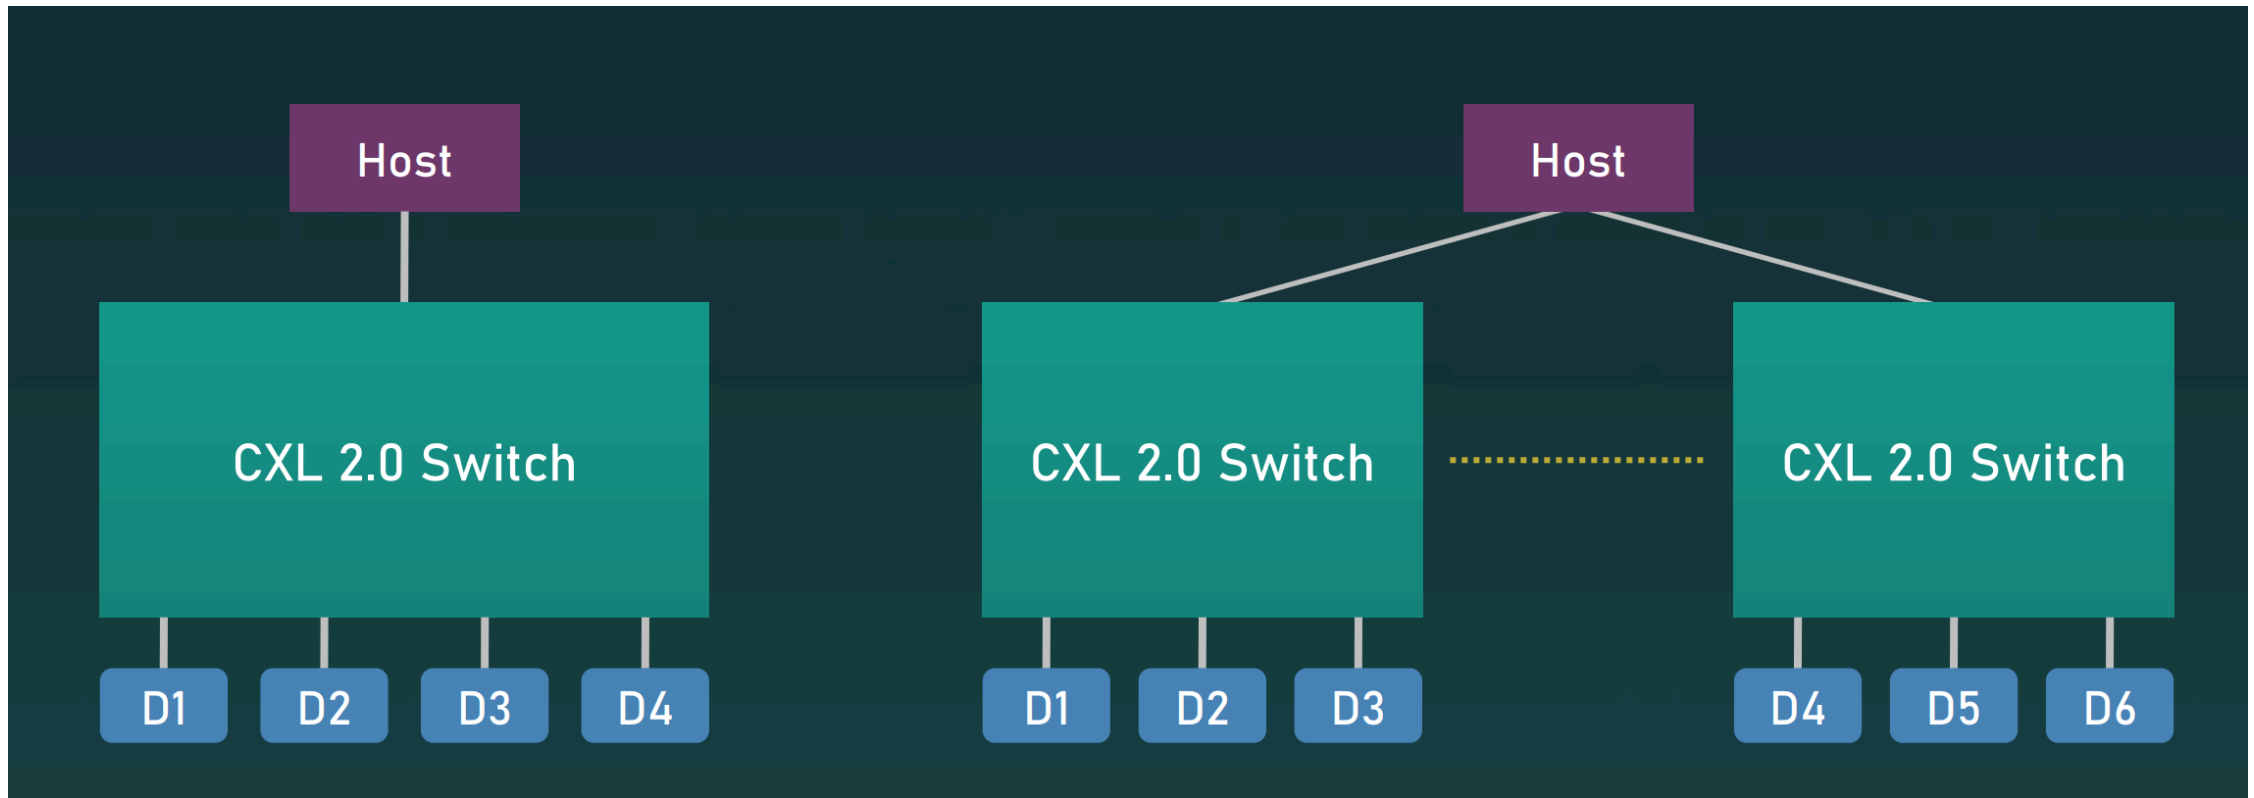
## CXL1.1□□□□

- CXL □□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□□: □□□□□□□□  
DDR3/4/5, LPDDR3/4/5, □□□□□□□□
- □□□□□□□□(persistence, latency, BW, endurance, etc)

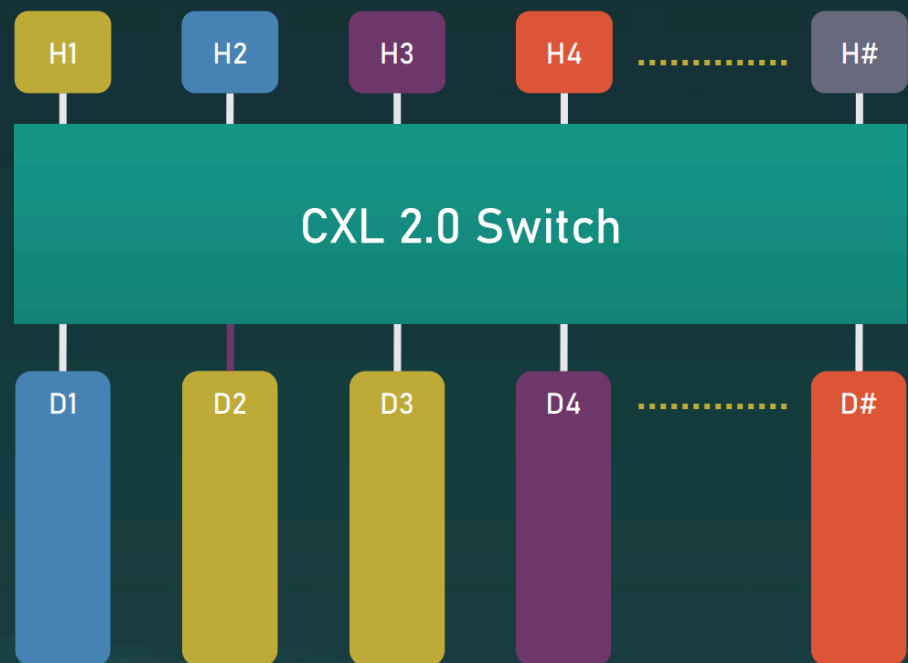


CXL2.0 switch

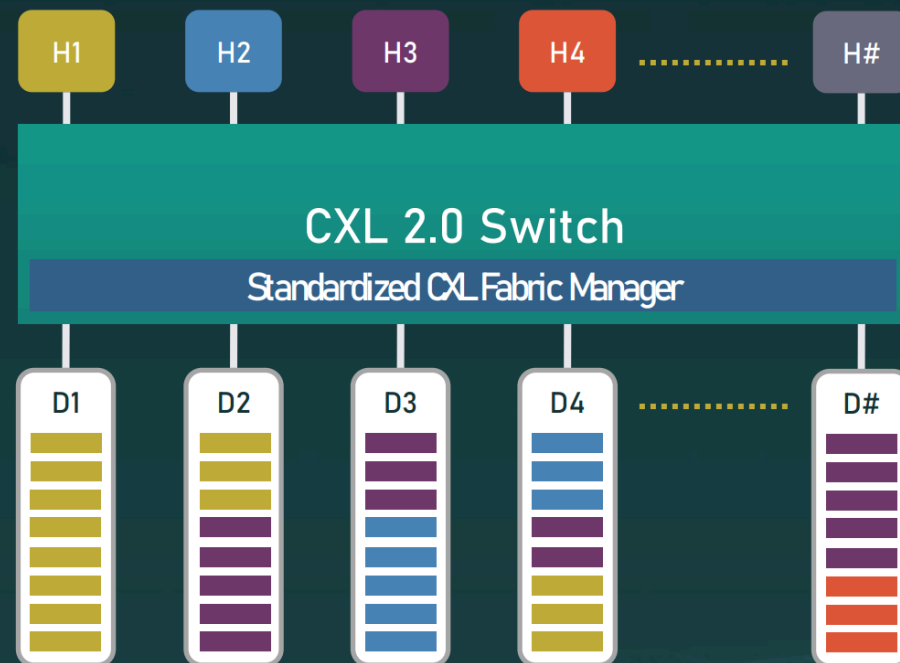
CXL2.0 □□single switch



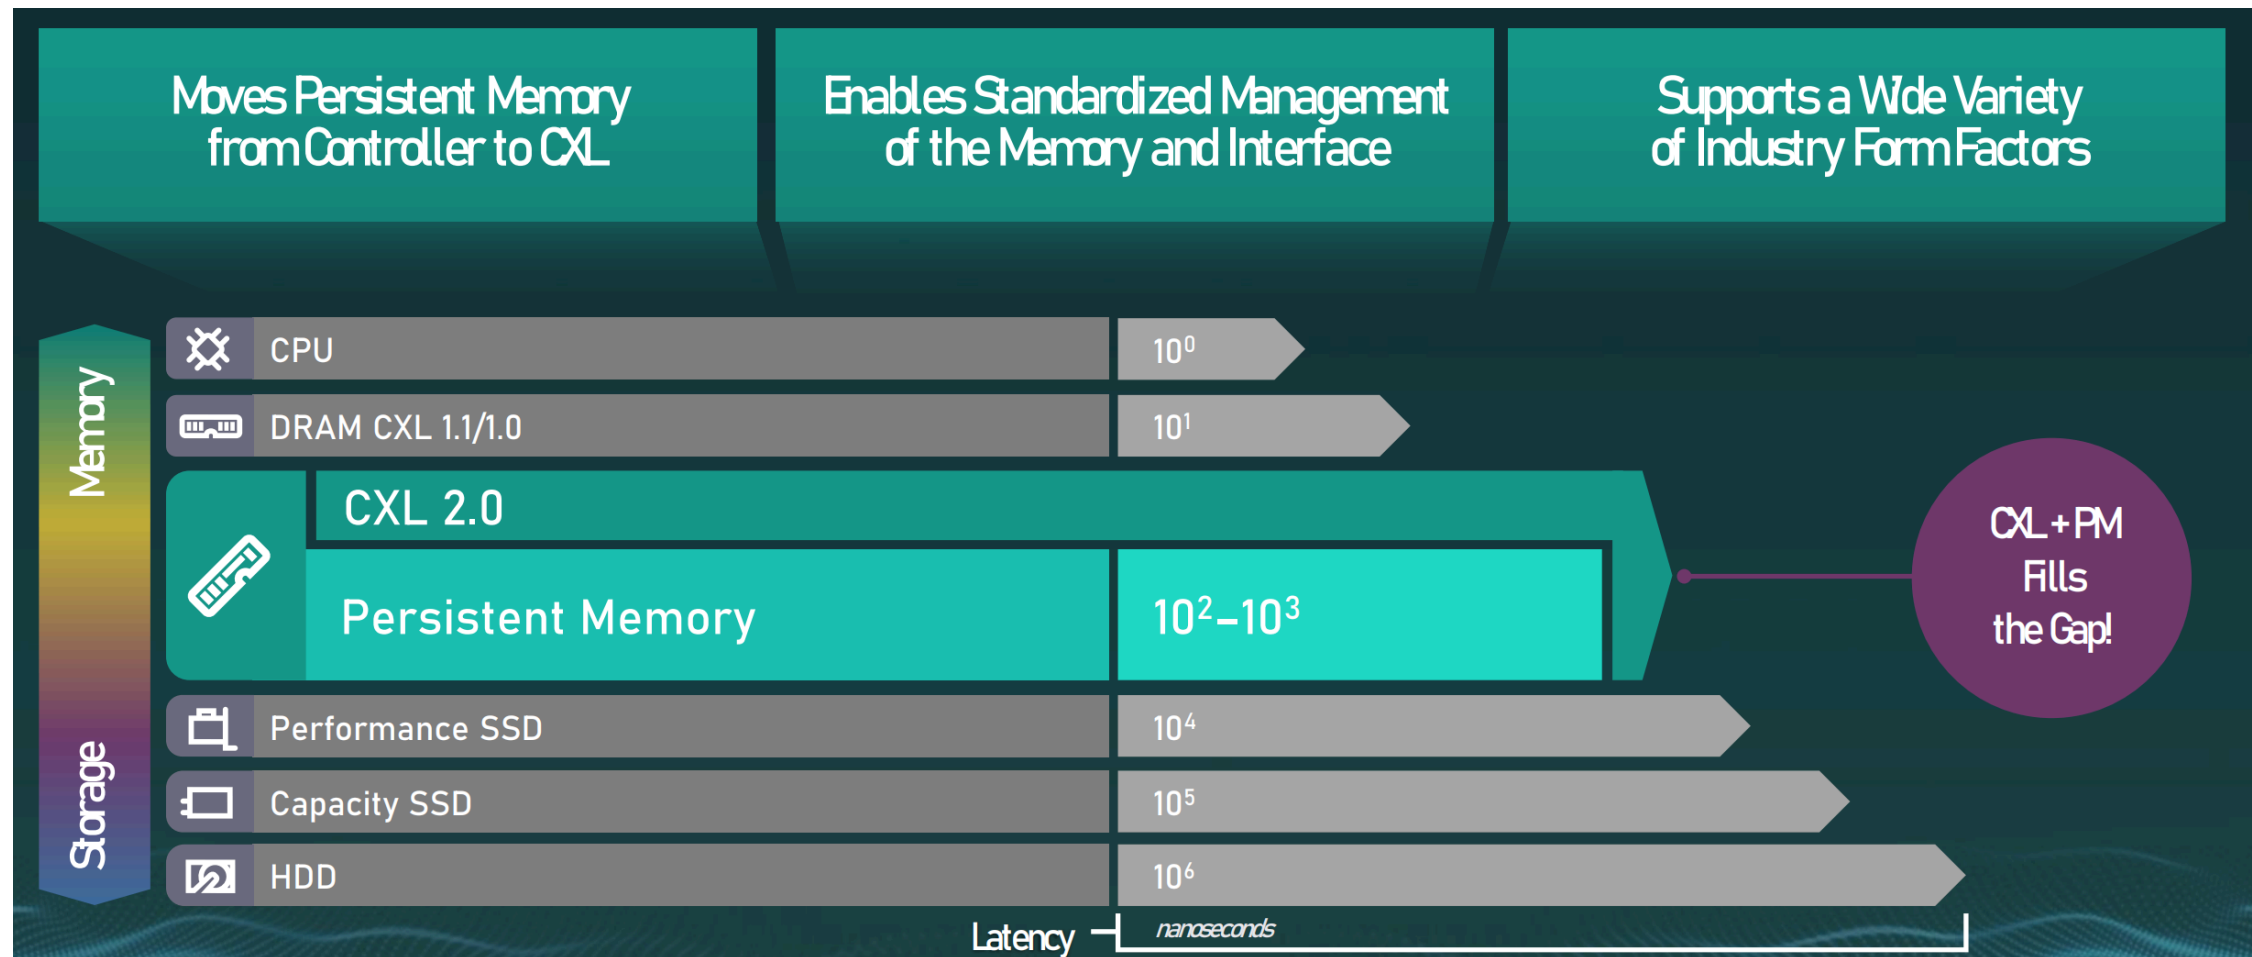
## Memory/Accelerator Pooling with Single Logical Devices



## Memory Pooling with Multiple Logical Devices



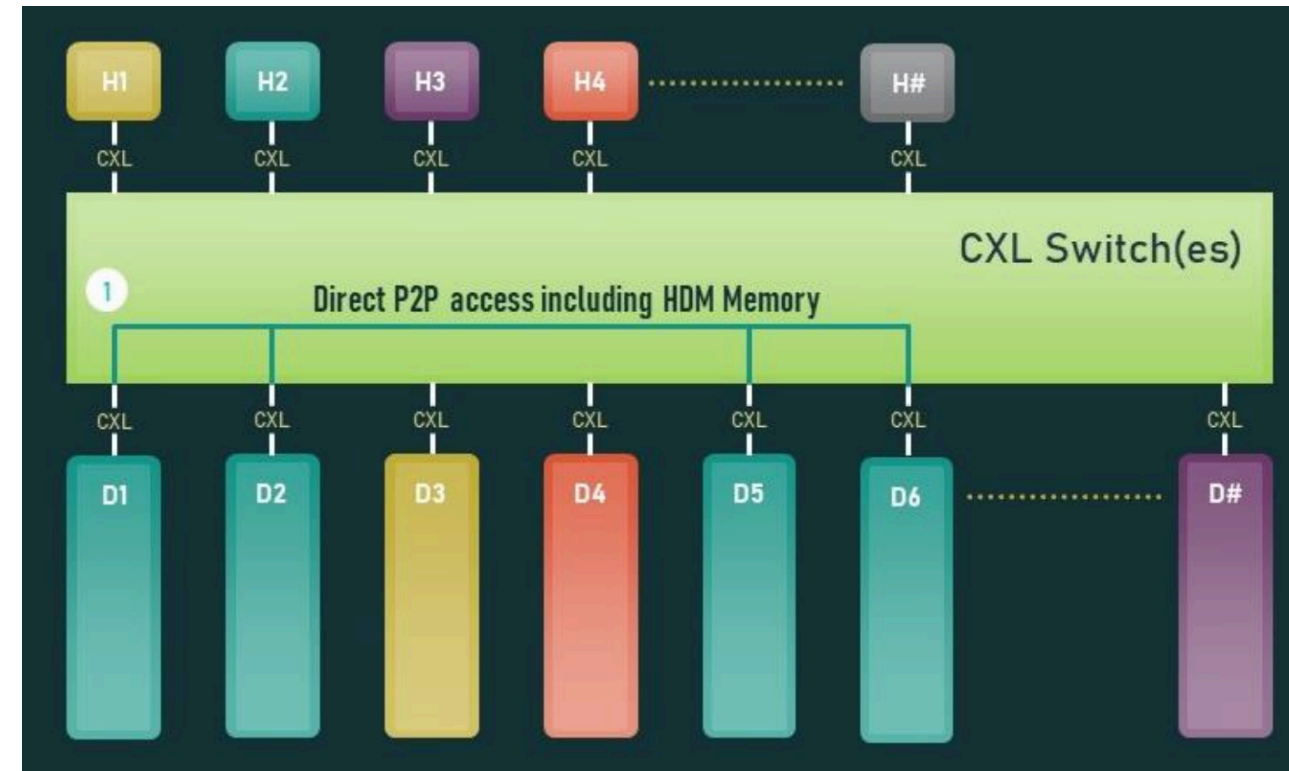
## CXL2.0 □□□□□



CXL3.0□□

CXL□□p2p□□□□□□□□

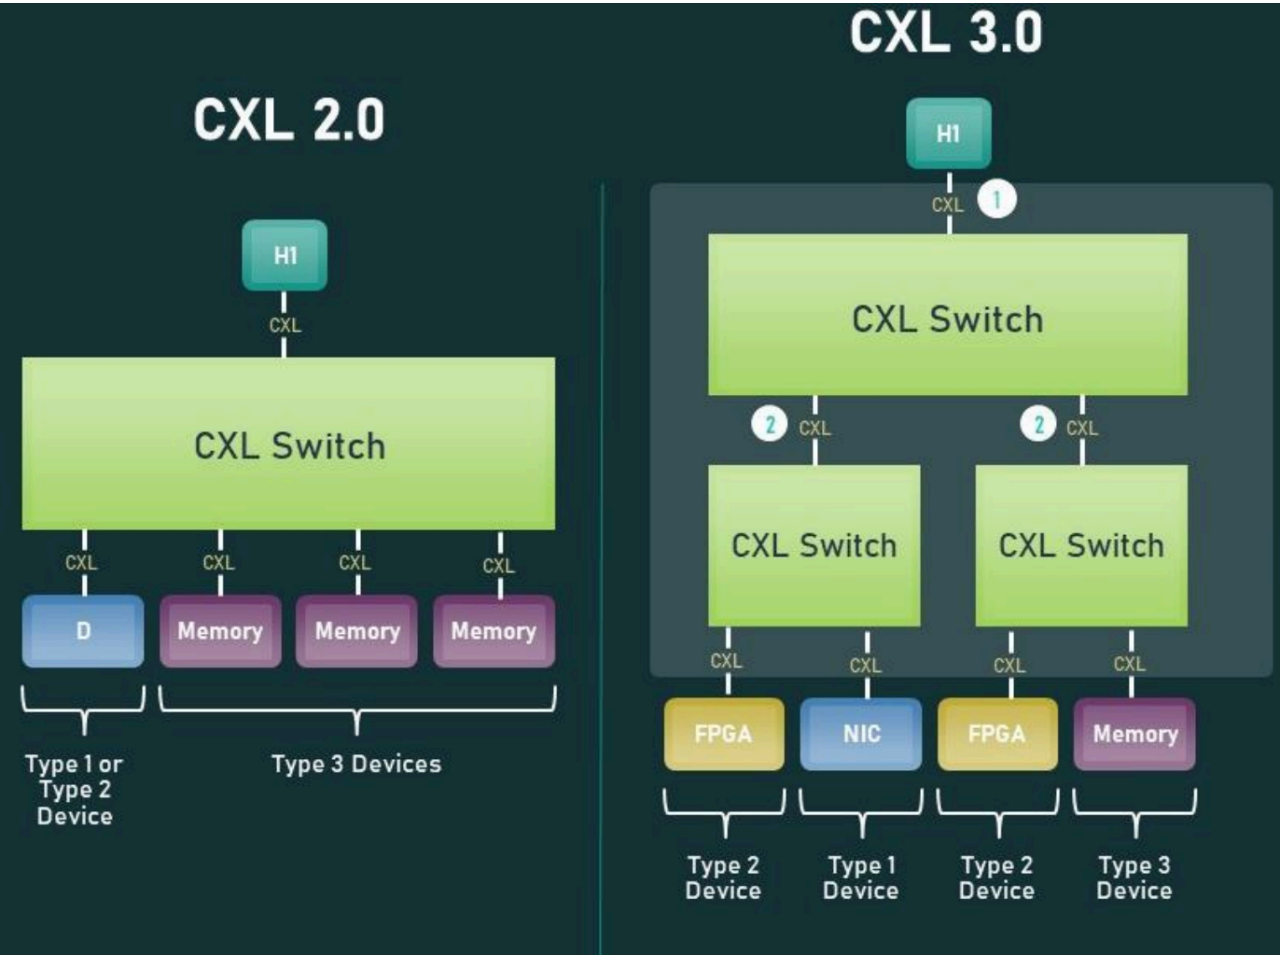
- Virtual hierarchies are associations of devices that maintains a coherency domain
- P2P to HDM-DB memory is I/O  
Coherent: a new Unordered I/O (UIO) Flow in CXL.io – the Type-2/3 device that hosts the memory will generate a new Back-Invalidation flow (CXL.Mem) to the host to ensure coherency if there is a coherency conflict



CXL3.0

16 CXL.cache

CXL

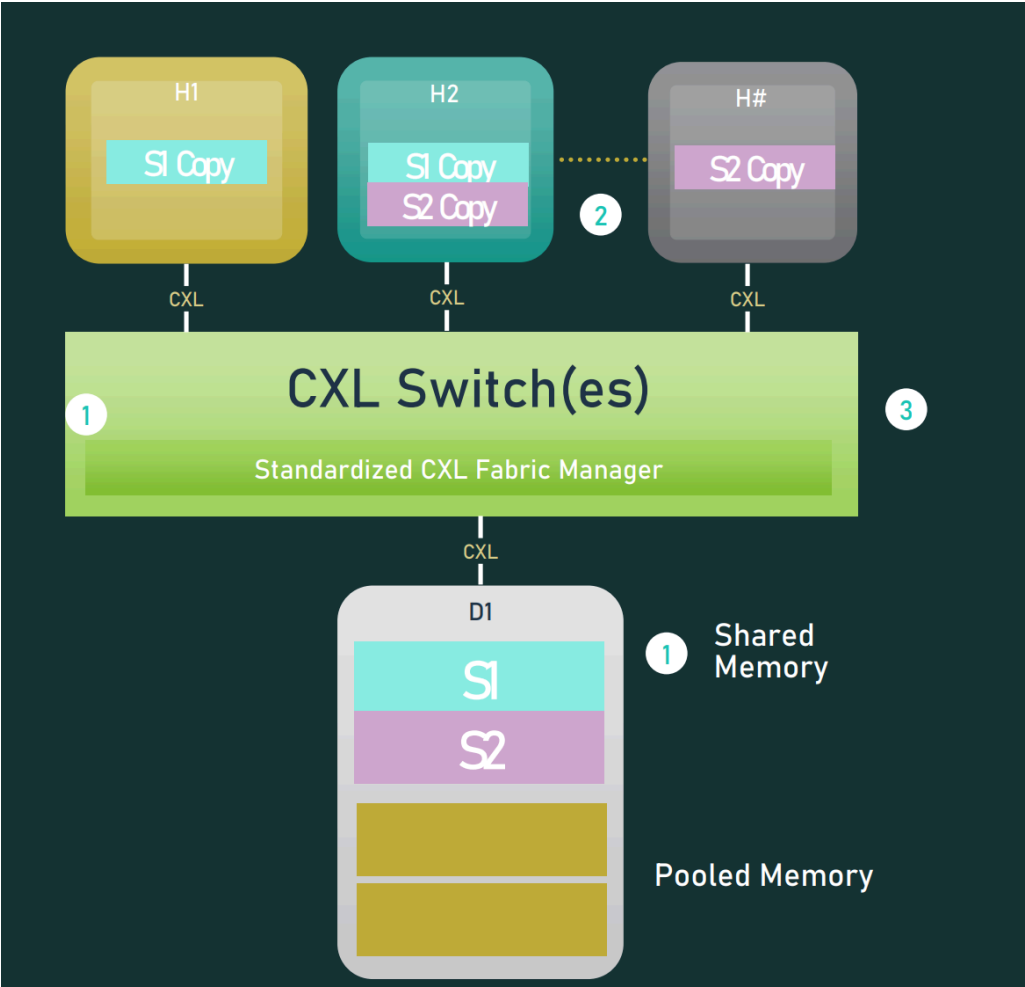


CXL3.0□□

□□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□□□  
□□□

**CXL 3.0** □□□□□□□□□□□□□□□□□□□□□□□□  
□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□  
□□□□

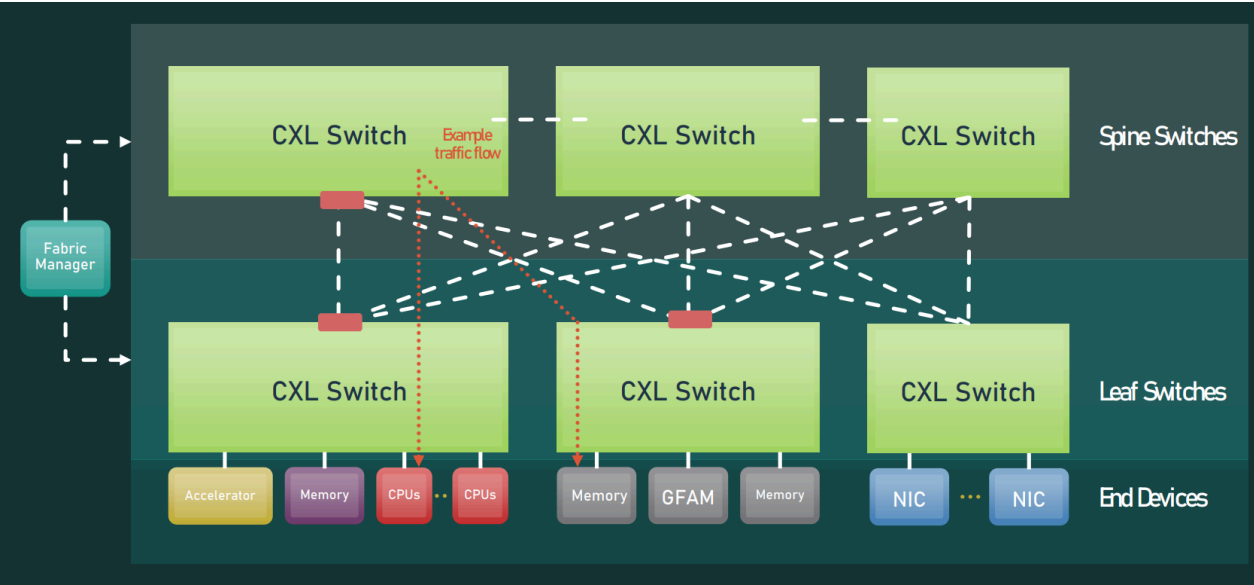




CXL3.0

CXL Fabric

CXL Fabric



## **CXL Fabric Improvements/Extensions**

- Scale out of CXL fabrics using PBR(Port Based Routing)

## **Trusted-Execution-Environment Security Protocol(TSP)**

- Allows for Virtualization-based, Trusted-Execution-Environments (TEEs) to host Confidential Computing workloads

## **Memory Expander Improvements**

- Up to 32-bit of meta-data and RAS capability enhancements

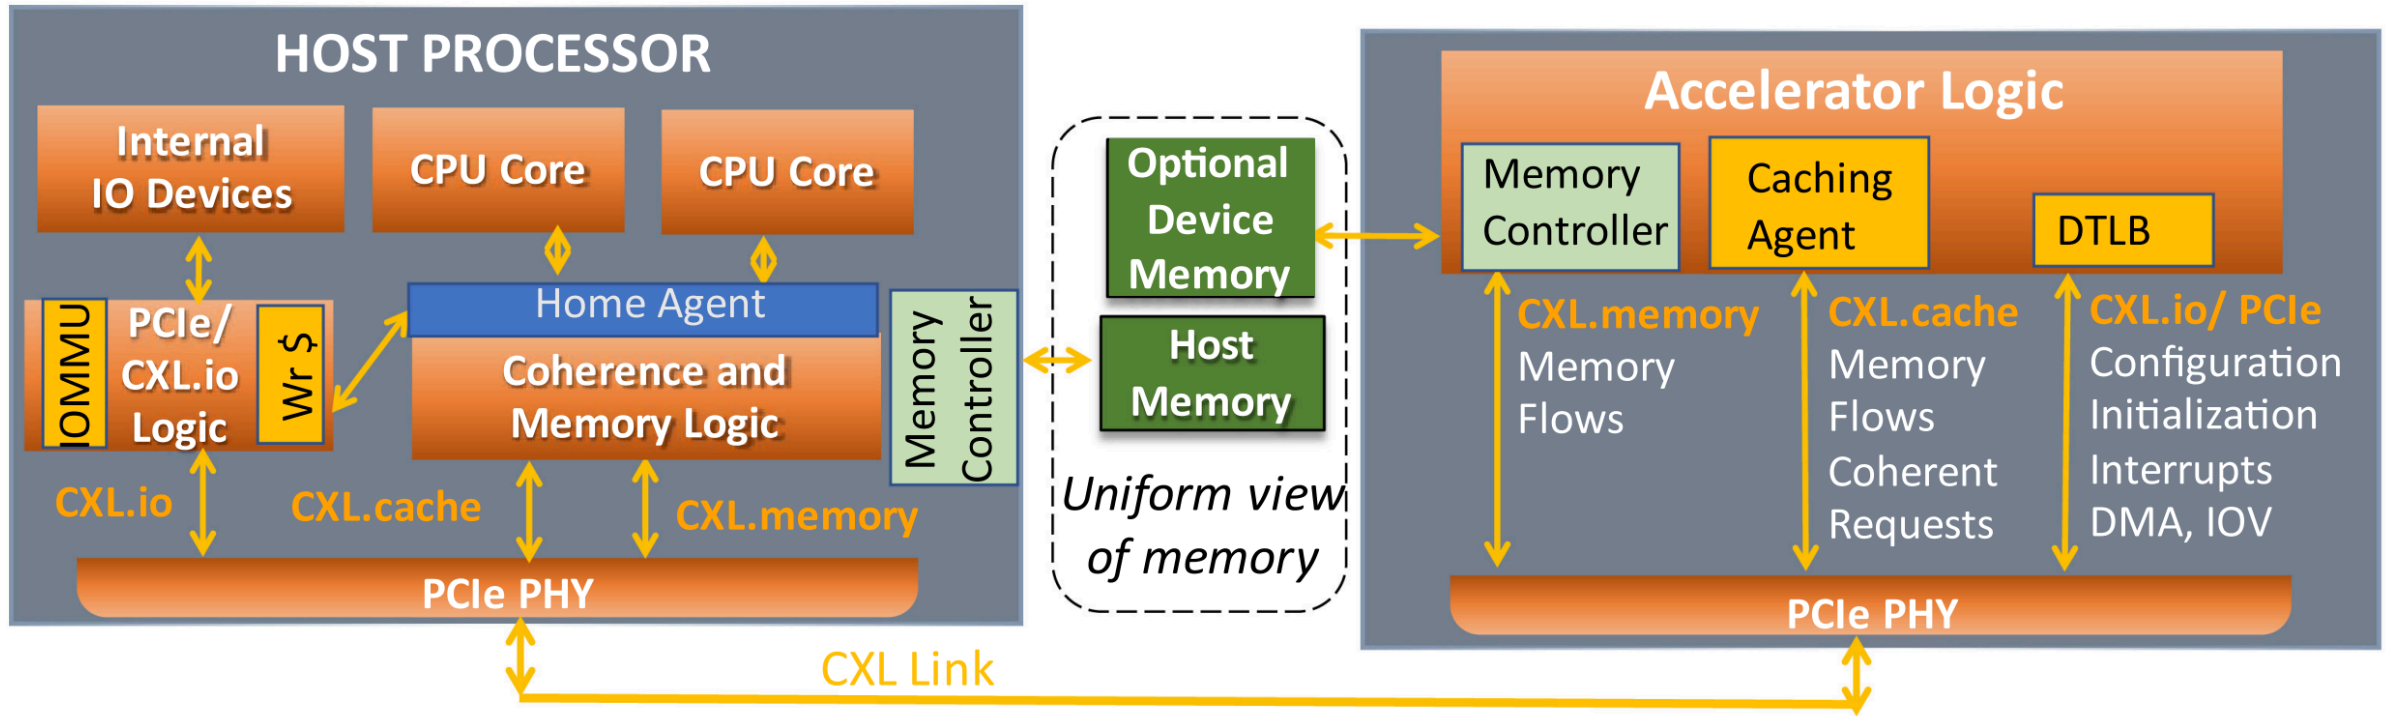
CXL ☐ ☐ ☐

CXL□□□□□□□□□□**CXL.io**□**CXL.cache**□**CXL.memory**□□□□PCle□□□□□□□□□□

CXL.io□□□□PCle□□□DMA□□

CXL.cache□□□□□□□□□□FPGA□□□□□□□□□□□□□□□□□□□□□□□□

**CXL.memory** Host-Managed Device Memory, HDM




## Type 1: Accelerator w/o Memory

### Usages:

- PGAS NIC
- NIC atomics

### Protocols:

- CXL.io
- CXL.cache




## Type 2: Accelerator w/ Memory

### Usages:

- GPU
- FPGA
- Dense
- Computation

### Protocols:

- CXL.io
- CXL.cache
- CXL.memory




## Type 3: Memory Buffer

### Usages:

- Memory BW expansion
- Memory capacity expansion
- Two-level memory

### Protocols:

- CXL.io
- CXL.mem



## CXL

CXL CPU CXL 1.1 CXL 2.0 CXL 3.0

Sapphire Rapids SPR CPU Agilex7 FPGA CXL AMD Genoa Bergamo CPU CXL SmartNIC ARM V2 N2 E2 CPU CXL 2.0

CXL1.1 ontage SK Hynix Microchip Micron Astera Labs CXL Type3 CXL1.1

CXL ☐ ☐ ☐ ☐ ☐

The diagram illustrates a memory layout with two rows of blocks. The top row contains 16 blocks, with the 4th block labeled 'CXL' and the 10th block labeled 'DRAM'. The bottom row contains 10 blocks, all of which are empty.

☐ ☐ ☐ ☐ ☐ ☐ ☐ ☐ CXL ☐ QoS ☐ ☐ ☐ ☐

☐☐☐☐☐☐☐☐☐☐☐☐☐☐ CXL ☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐ UCle ☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐☐  
☐☐☐☐☐

# 汇报完毕，恳请指正

宋敬炎